



Deep Learning

Content


- **Vanishing Gradient & Activation Functions**
- **Dropout**
- **Batch Normalization**



Gradient Vanishing & Activation Functions

Gradient Vanishing & Exploding

- **Gradient is easy to vanish or explode**
 - To many terms are multiplied.
 - If some are small numbers, gradient becomes very small.
 - If some are large numbers, gradient becomes very large.



$$\frac{\partial E_n}{\partial w_{kj}} = -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) h_{nj}$$



$$\frac{\partial E_n}{\partial w_{ji}} = -h_{nj} (1 - h_{nj}) x_{ni} \sum_{k=1}^m w_{kj} (t_{nk} - o_{nk}) o_{nk} (1 - o_{nk})$$

$$\frac{\partial E}{\partial w_{ip}} = \left(\sum_{j=1}^J \left(\sum_{k=1}^K -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) w_{kj} \right) h_{nj} (1 - h_{nj}) w_{ji} \right) h_{ni} (1 - h_{ni}) h_{np}$$

Activation Function

■ Vanishing Gradient

- The major terms are the derivatives of the activation function



$$\frac{\partial \text{Sigmoid}}{\partial w} \leq \frac{1}{4}$$


$$\frac{\partial E_n}{\partial w_{kj}} = -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) h_{nj}$$

$$\frac{\partial E_n}{\partial w_{ji}} = -h_{nj} (1 - h_{nj}) x_{ni} \sum_{k=1}^m w_{kj} (t_{nk} - o_{nk}) o_{nk} (1 - o_{nk})$$


$$\frac{\partial E}{\partial w_{ip}} = \left(\sum_{j=1}^J \left(\sum_{k=1}^K -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) w_{kj} \right) h_{nj} (1 - h_{nj}) w_{ji} \right) h_{ni} (1 - h_{ni}) h_{np}$$

Activation Function

- Using another functions instead of sigmoid
 - Rectified Linear Unit (ReLU)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$



$\frac{\partial E_n}{\partial w_{ji}}$ = Some formular with 3 mulitiplications of $\frac{\partial f}{\partial w}$

$\frac{\partial E_n}{\partial w_{hg}}$ = Some formular with 4 mulitiplications of $\frac{\partial f}{\partial w}$

Activation Function

■ Advantage

- No vanishing gradient problems.
 - Deep networks can be trained without pre-training
- Sparse activation
 - In a randomly initialized network, only about 50% of hidden units are activated
- Fast computation:
 - 6 times faster than sigmoid function

■ Disadvantage


- Knockout Problem

Activation Function

- You may use another


- Leaky ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$




- Swish (or SiLU-Sigmoid Linear Unit)

$$f(x) = \frac{x}{1 + e^{-x}}$$



Other Activation Functions



Activation Function

■ Summary


- Sigmoid functions and their combinations generally work better but are sometimes avoided due to the vanishing gradient problem
- ReLU function is a general activation function and is used in most cases these days
- If we encounter a case of dead neurons in our networks the leaky ReLU function is the best choice
- ReLU function is usually used in the hidden layers
- As a rule of thumb, you can begin with using ReLU function and then move over to other activation functions in case ReLU doesn't provide with optimum results





Regularization

Overfitting


■ Overfitting



Underfit



Overfit




Regularization

- **What is Regularization**
 - Introducing additional information to prevent over-fitting
- **Approaches**
 - Proper Learning: Early stopping
 - Proper Structure: Weight decay, Dropout, DropConnect, Stochastic pooling

Early Stopping

- Split data into 3 groups



Weight Decay

■ L1 Regularization

- Leading most weights very close to zero
- Choosing a small subset of most important inputs
- Resistant to noise in the inputs.

$$\tilde{E}(\mathbf{w}) = E(\mathbf{w}) + \frac{\lambda}{2} |\mathbf{w}|$$


■ L2 Regularization

- Penalizing peaky weights
- Encouraging to use all of its inputs a little rather than using only some of its inputs a lot.


$$\tilde{E}(\mathbf{w}) = E(\mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

Weight Decay


- **Complex Structure vs Simple Structure**



Node
Pruning




Link
Pruning




Weight Decay

- **Complex Structure vs Simple Structure**



Set
many links
to zero




$|w|$ is large \leftrightarrow NN is Complex

$|w|$ is small \leftrightarrow NN is Simple

Weight Decay

- **Example: Separating green and red**




L2 regularization strengths of 0.01, 0.1, and 1

Dropout

■ In a complex Neural Network

- All nodes do not take the same amount of responsibility
 - While training, some nodes are correlated
- All nodes are not equally trained
 - Some nodes are trained much, but some are not





- If the output of the node is bad, the connection weight will decrease.
- If connection weight is close to 0, precedent connection weights are hardly trained.

Dropout

- **How can we reduce the structural complexity?**


- Let's simply remove some nodes, and
- Train the simplified neural network
- Hmm??




Dropout

- **Do this at every epoch**


- Randomly choose nodes with a probability of p
 - Usually $p = 0.5$
- Train the simplified neural network
 - At every epoch, we train different neural network which share connection weight each other




Original
network



Update 1



Update 2




Update 3

...

Dropout

- **Testing**
 - Use all the nodes without dropout




Dropout

■ Testing


Training of Dropout

Assume dropout rate is 50%



Testing of Dropout


No dropout



Dropout

- **The effect of the dropout rate p :**

- An architecture of 784-2048-2048-2048-10 is used on the MNIST dataset.
- The dropout rate p is changed from small numbers (most units are dropped out) to 1.0 (no dropout).




Dropout

- **The effect of data set size:**

- An architecture of 784-1024-1024-2048-10 is used on the MNIST dataset.

Extremely small data set
- Dropout does not improve error rate, and even makes it worse.

Average to large data set
- Dropout improves error rate.



Huge data set
- Dropout barely improves the error rate. The data set is big enough, so that overfitting is not an issue.

Dropout

■ Summary


- Dropout is a very good and fast regularization method.
- Dropout is a bit slow to train (2-3 times slower than without dropout).
- If the amount of data is average-large – dropout excels. When data is big enough, dropout does not help much.
- Dropout achieves better results than former used regularization methods (Weight Decay).

Batch Normalization

Batch Normalization

- **Covariate Shift**

- A change in the distribution of a function's domain.





- Can your model work properly?


Batch Normalization

■ Internal Covariate Shift

- Input distribution of the red node



- While learning, red connection weights will change based on the input distribution
- After learning, the whole connection weights changes, which cause the change of the input distribution
- The assumption of the learning is broken




Batch Normalization

- **Internal Covariate Shift**

- It disturbs the learning process,
- Learning is getting slow down


- **What shall we do?**

- Why don't we normalize the distribution of inputs





Batch Normalization


Input Normalization



$$\mu, \sigma^2 \quad \hat{net} = \frac{net - \mu}{\sqrt{\sigma^2 + \epsilon}}$$



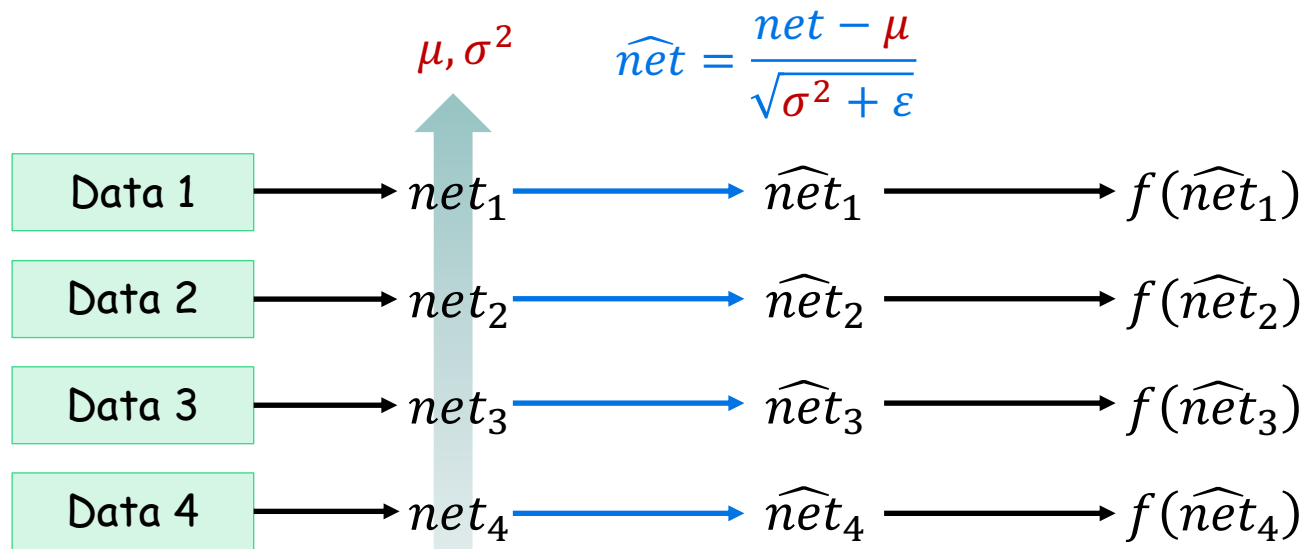

Distorted distribution



Normalized distribution

Batch Normalization



- Input Normalization



?

Batch Normalization

Input Normalization



$$net = \mathbf{w}\mathbf{x} + b$$


$$E(net) = E(\mathbf{w}\mathbf{x}) + b$$

$$\widehat{net} = \frac{\mathbf{w}\mathbf{x} + b - (E(\mathbf{w}\mathbf{x}) + b)}{\sqrt{\sigma^2 + \epsilon}}$$

$$\widehat{net} = \frac{\mathbf{w}\mathbf{x} - E(\mathbf{w}\mathbf{x})}{\sqrt{\sigma^2 + \epsilon}}$$

Batch Normalization


Input Normalization



μ, σ^2


$$\widehat{net} = \frac{net - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$\widetilde{net} = \gamma \widehat{net} + \beta$$



Batch Normalization


Input Normalization



μ, σ^2



$$\widehat{net} = \frac{net - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$\widetilde{net} = \gamma \widehat{net} + \beta$$



Batch Normalization

- For a Single Node




$$\mu_B = \frac{1}{D_B} \sum_{d=1}^{D_B} net_{di}$$

$$\sigma_B^2 = \frac{1}{D_B} \sum_{d=1}^{D_B} (net_{di} - \mu)^2$$

Batch Normalization

■ Testing

- For Training, the mean and variance of each batch are used for normalization
- For Testing, of which data the mean and variance will be used?
 - Estimated with those of batches in the training




Batch Normalization

■ Advantage


- Reduces internal covariant shift.
- Reduces the dependence of gradients on the scale of the connection weights.
- Regularizes the model and reduces the need for regularization techniques.
 - It adds some stochastic noise to the activations as a result of using noisy estimates computed on the mini-batches. This has a regularization effect in some applications,

Batch Normalization


- Performance with BN



(a) *accuracy*



(b) Without BN



(c) With BN

input distributions

Batch Normalization

■ Disadvantage

- Expensive: Memory and time
 - Must keep interim results of all instances in a batch
 - Especially in CNN, usually an image is large
- Hard to apply when the batch size is small
 - If batches are small, the means and variances cannot approximate the global ones.
- Hard to apply to recurrent networks
 - It doesn't match to structure of recurrent networks
 - Hard to implement with recurrent networks

Layer Normalization

- **Recap: Batch Normalization**
 - Normalization of each node output

Batch normalization

$$\begin{aligned}\mu_j &= \frac{1}{m} \sum_{i=1}^m x_{ij} \\ \sigma_j^2 &= \frac{1}{m} \sum_{i=1}^m (x_{ij} - \mu_j)^2 \\ \hat{x}_{ij} &= \frac{x_{ij} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}\end{aligned}$$

i, j : index of the batch and the node of hidden layers

Layer Normalization

■ Recap: Batch Normalization

- Normalization of each node output

Outputs of a layer for a mini-batch					Normalized outputs		
			Avg	Std			
1.0	3.0	6.0	3.3	2.5	-0.9	-0.1	1.1
2.0	2.0	2.0	2.0	0.0	0.0	0.0	0.0
0.0	1.0	5.0	2.0	2.6	-0.8	-0.4	1.1
4.0	6.0	1.0	3.7	2.5	0.1	0.9	-1.1
5.0	2.0	3.0	3.3	1.5	1.1	-0.9	-0.2
1.0	0.0	1.0	0.7	0.6	0.6	-1.2	0.6

Layer Normalization


- **Proposed as an alternative to Batch Normalization**
 - Works regardless of batch size (batch size = 1)
 - Performs well with RNNs

Layer normalization:

$$\begin{aligned}\mu_i &= \frac{1}{m} \sum_{j=1}^m x_{ij} \\ \sigma_i^2 &= \frac{1}{m} \sum_{j=1}^m (x_{ij} - \mu_i)^2 \\ \hat{x}_{ij} &= \frac{x_{ij} - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}\end{aligned}$$

i, j : index of the batch and the node of hidden layers

Layer Normalization



Layer Normalization

- **Group Normalization shows consistent accuracy with smaller batches**
 - Tested on ImageNet (1000 Classes, 1.28M training, 50K validation), ResNet-50

